

Lab Report

Mars Lander

A guide to the modifications to the original source code to complete core tasks of developing an autopilot for landing and algorithms for numeric integration. In addition, notes and theory to complete some of the additional exercises.

Oliver Jones
Christ's College Lab Group 115
26.09.2017

1. INTRODUCTION

The purpose of this task is to introduce the C++ programming language as well as an introduction to basic control theory. The importance of the choice of integration algorithm is key as the instabilities introduced can cause unpredictable behaviour. This is initially tested in a simple 1-dimensional spring system, before being implemented in 3D.

2. SPRING SIMULATION

2.1 INTEGRAL TYPES

The first task is to complete a Verlet integration algorithm for the provided Spring source code. The initial code demonstrates Euler integration and it can be shown that a smaller time-step will reduce the instability in the system, however it cannot be removed completely.

$$x_{n+1} \approx 2x_n - x_{n-1} + \Delta t \frac{d^2 x_n}{dt^2} \quad [1]$$

The most important part of dealing with Verlet integrals is the initial calculation. As shown in [1], the current position and the previous position is required to calculate the next time-step. There are two possible ways of doing this:

- Apply the first time-step as a complete Euler integral and saving this for the next iteration
- Initialise the x_{n-1} variable as $x_{n-1} = x_n - \Delta t \frac{dx_n}{dt}$

Both options produce similar results though for the purposes of simpler code, it is faster to initialise x_{n-1} as an appropriate value then use Euler, as no acceleration or new velocity is required to be calculated.

2.2 CONVERSION TO C++

Processing the simulation through C++ has the obvious advantages of speed. With compiler optimisation, the code can run up to 20x faster than its Python equivalent. In an effort to further understand data manipulation in C++ additions have also been made that rather than writing the code out to a .txt file, each array is saved as its own binary file which can be read into python in a similar manner. Such changes do have a marginal effect on the performance as no ASCII conversion is required.

Similar additions are the use of command line arguments to choose the integrator type, the functions for which are kept in a separate file to better understand headers and forward declaration. If no such arguments are supplied the application requests a choice from the user.

Finally, the source code was modified to automatically call the python function and generate the necessary plots to show the evolution of the system over time.

3. ORBITAL SIMULATIONS

Extending the previous task to 3D is fairly trivial when using 3 element arrays. In python, setting up some simple vector functions is straight forward particularly with numpy arrays. Setting the origin as the planetary centre, a circular orbit can be established by setting the velocity of the lander at 90 degrees to the position vector with magnitude as follows:

$$v_0 = \frac{GM}{r_0} \quad [2]$$

Using Newton's law of gravity, the velocity can then be updated as necessary to predict the path of the orbit. Multiplying v_0 by a scalar can allow simple conic paths to be shown

Orbits follow paths defined by conics – the eccentricity can be used to determine how stable an orbit will be. $e < 1$ will be a stable ellipse, reducing to perfectly circular at $e = 0$, $e > 1$ gives a parabolic orbit at the orbiter will shoot off into space and not be trapped by the planet's gravity, $e < 0$ give chaotic hyperbolic orbits.

4. SIMULATOR CORE TASKS

4.1 NUMERICAL DYNAMICS

The original source code, when compiled, merely shows a stationary lander that does not move in any of the scenarios, as the code does not have any dynamics calculations in it. Unlike the Orbits calculations done in Python, drag is also now a factor, as well as the thrust from the rockets of the lander. Rocket thrust is already provided, and just requires addition into the dynamics but all other calculations must be created.

To keep the code tidy a 'dynamics' file was created. Initially this contained just the additional functions that calculate forces on the lander, but in later extension exercises more functions are added in to keep act as a single location to keep track of all additional functions not in the original source code.

The original functions in 'dynamics' is the calculation of gravity based on the lander's position, and separate lander and parachute drag calculations using the following equation:

$$\mathbf{F}_d = -\frac{1}{2}\rho C_d A |\mathbf{v}|^2 \mathbf{v} \quad [3]$$

Atmospheric density is also provided in a function, which follows an exponential decay as the absolute position increases and cuts off completely at the edge of the Exosphere (altitude of 200km). The total area of the lander is 10x the lander radius, and for the purposes of simplicity all drag calculations assume the lander is moving with the base pointing in the direction of travel.

To deal with the integrators, a similar method is described as above in section 2. However, as a static variable is declared for 'old_position', the value cannot be initialised straight away, as when the scenario is changed, it is not reset. Instead, a control statement sets its value when the simulation time is 0 (which it is reset to when the scenario changes).

4.2 AUTOPILOT

Once the numerical dynamics have been sorted correctly, the lander can be controlled manually. The difficulty in landing by hand in any of the scenarios, even with parachute assistance, is difficult, and an autopilot is required to consistently land safely. To land safely, the lander must not be moving more than 1m/s in any direction.

The autopilot will be based on a proportional control system, and the descent rate should reduce linearly as the surface approaches:

$$\mathbf{v} \cdot \mathbf{e}_r = -(\mathbf{v} \cdot \mathbf{e}_r)_0 - K_h |\mathbf{r}| \quad [4]$$

Where the ideal value of $\mathbf{v} \cdot \mathbf{e}_r$ is the descent rate we want to be at when the lander touches the ground, and K_h is a positive constant. An error term can be defined as the difference between the two sides of the equation:

$$\varepsilon = -(\mathbf{v} \cdot \mathbf{e}_r)_0 - K_h |\mathbf{r}| - \mathbf{v} \cdot \mathbf{e}_r \quad [5]$$

This gives us our Proportional output:

$$P_{\text{out}} = K_p \varepsilon \quad [6]$$

Again K_p is a positive constant. At our ideal velocity, there should be no acceleration so the thrust from the lander should balance its weight:

$$\Delta = \frac{GMm}{r^2 T_{\text{max}}} \quad [7]$$

Where T_{max} is the maximum thrust capable from the rockets. With these definitions, we can define the throttle regime:

$$\text{throttle} = \begin{cases} 0 & P_{\text{out}} \leq -\Delta \\ \Delta + P_{\text{out}} & -\Delta \leq P_{\text{out}} < 1 - \Delta \\ 1 & P_{\text{out}} \geq 1 - \Delta \end{cases} \quad [8]$$

For information on selecting the values of K_h refer to section 5.4.

Determining Parachute Release

The simulator will cause the parachute to be lost if the force on it exceeds a certain force (causing tethers to snap) or if the lander is travelling too fast (causing it to vaporise). Therefore, it is important to have a logical system that will decide an appropriate time to release the parachute to make the most use of its breaking abilities while not causing it to be lost.

The source code provides an indicator for when either of these two conditions are met, and initial systems simply checked whether these values are true or false. This works well for scenario 1 as the parachute can be released instantly with no issues, however in scenario 5, as the lander is at the edge of the exosphere, there is not sufficient atmosphere to stop the lander accelerating, causing the parachute to vaporise.

It is safe to assume that when the thrusters fire the lander will not be exceeding the maximum safe conditions once it returns from that regime. Therefore, the system was modified with a second Boolean to confirm whether the initial firing of the rockets had begun, if so and safe to deploy parachute it will. This provides a much more reliable way of deploying the parachute, though has issues in scenario 1 in that it can be released much earlier than when the thrusters are first fired.

Originally, to cover all possible scenarios, a prediction function was created that will return true if it is *useful* for the parachute to be released now. This used a series of virtual parameters, and predicted how the velocity of the lander will evolve over time if the parachute is released now¹. The function would return false if it predicted evolution into an unsafe regime or if the lander virtually crashes, but returns true if the velocity reduces before any of these happen. This provides an adaptive parachute launcher which, when combined with the autopilot, allows the lander to touch down successfully in all the scenarios that don't have stable orbits.

However, it was later realised that if the altitude is less than 50km then the parachute will be useful if it is also safe to deploy so the above function was no longer used to reduce processing, though would provide a more dynamic calculation if planetary parameters were changes such as atmospheric density or size.




Figure 1a Scenario 5 Error and Velocity vs Altitude Fuel Efficient Landing




Figure 1b Scenario 5 Error and Velocity vs Altitude Fuel Soft Landing

¹ To save computing power, the function used Euler integrals with a slightly larger time step than the actual simulation as only a rough estimate is required for reliable results.

5. EXTENSION TASKS

5.1 AEROSTATIONARY ORBIT

To produce an aerostationary orbit, the lander must be moving such that its orbital period is equal to that of Mars. From Kepler's law of planetary motion:

$$T^2 = \frac{4\pi^2}{GM} r^3 \quad [9]$$

Setting T as appropriate allows the calculation of the appropriate radius. Then from simple circular motion, plugging this into [2] means that we can calculate the appropriate speed to maintain that circular motion.

5.2 ANY ANGLE ATTITUDE CONTROL

Currently the attitude stabilisation set the orientation of the lander such that its base is facing the centre of Mars. To allow for planetary rotation and wind, any angle attitude control must be possible. Fortunately, we already have a set of vectors aligned with the centre of the lander as well. By crossing one of these vectors with the position, we get an axis that is perpendicular to the plane of the screen which we can rotate about.

To form a rotation about an arbitrary axis we can use the following rotation matrix:

$$R = \begin{pmatrix} \cos \theta + u_x^2(1 - \cos \theta) & u_x u_y(1 - \cos \theta) - u_z \sin \theta & u_z u_x(1 - \cos \theta) + u_y \sin \theta \\ u_x u_y(1 - \cos \theta) - u_z \sin \theta & \cos \theta + u_y^2(1 - \cos \theta) & u_z u_y(1 - \cos \theta) - u_x \sin \theta \\ u_z u_x(1 - \cos \theta) + u_y \sin \theta & u_z u_y(1 - \cos \theta) - u_x \sin \theta & \cos \theta + u_z^2(1 - \cos \theta) \end{pmatrix} \quad [10]$$

Where u is the normalised axis of rotation and θ is the angle of rotation about that axis. Applying this matrix does allow for some rotation control, however due to a phenomenon called gimble lock, getting into a certain orientation then removes one degree of freedom. Furthermore, calculating the elements of the matrix is computationally expensive and can slow down the programme, particularly at higher speeds.

An alternative to the rotation matrix is using a technique called quaternions. This is a set of complex numbers extended to a 4-dimensional space (3 imaginary and one real). Many computer games use this technique over matrices as they do not have the issues described above.

$$P = \cos \frac{\theta}{2} + (u_x \mathbf{i} + u_y \mathbf{j} + u_z \mathbf{k}) \sin \frac{\theta}{2} \quad [11]$$
$$\mathbf{v}' = P \mathbf{v} P^{-1}$$

Applying the rotation quaternion P as defined above on the 'up' vector \mathbf{v} (converted to a quaternion with a real part of 0), and will rotate the given 'up' vector in the attitude stabilisation function to the desired orientation, causing the lander to turn. Fortunately, many quaternion functions had already been defined in the original source code so this could be implemented in a few lines. An additional

statement was also added that if the stabilised attitude angle was approximately 0, then set up to the normalised position vector to save the computation.

Additional controls were then also added to allow manual changing of the angle by incrementing the stabilised attitude angle.

5.3 PLANETARY ROTATION MECHANICS AND WIND

Rotation

Currently, while the orbital view visuals show the rotation of the planet, the mechanics themselves do not take this into account. To correct this, a new relative velocity variable needs to be defined. This is calculated by using similar methods as described in section 5.1, by saying 'if the lander was on the surface now, how fast would it need to be moving to stay on the same spot on the surface?' By adding this additional speed onto ground speed and drag calculations, the rotation of the atmosphere will accelerate the lander in the appropriate direction. Furthermore, when the lander has 0 velocity it will now have a significant speed in relation to the surface.

Wind

Adding wind mechanics to this was straightforward. A new Dynamics function was created to create a normal distribution about a mean wind value and return a wind speed to the drag calculations. Using a normal distribution allows for slight variations in the wind, and increasing the standard deviation of the distribution increases how gusty and varying the wind can be.

A new switch was added to the 'w' key to allow the wind to be toggled on and off.

Modifying Autopilot to cope

With planetary rotation, there is now a lateral motion that needs to be considered. Using the parachute, no change needs to be made as the drag provides the necessary force to reduce the ground speed to 0. However, without the parachute, thrusters now need to fire at an angle to slow both descent and ground speed.

Initial solutions were to set the stabilised attitude angle such that the lander pointed against the direction of the velocity. However, this caused an odd bug where, just as the lander was about to touch down, the lander would flip on its side causing the ground speed to suddenly increase and crash. The second solution was to set the 'up' vector in the stabilised attitude function to the direction of velocity, but again this caused similar bugs. It was later realised that the 'thrust_wrt_world' function has a simplifying control statement where it assumes the lander is pointing directly up if 'stabilized_attitude_angle' is 0 to simplify processing which caused the bug.

When the wind speed increases there is the danger that the parachutes drag will work against the thrusters so the autopilot will eject the parachute when the ground speed of the lander approaches the wind speed and it is near the ground, so that the wind does not drag the lander into a crash.

5.4 TUNING AUTOPILOT

Tuning the values of K_p and K_h allows you to produce an autopilot regime that can either minimise fuel usage or force on the lander. The least fuel possible would mean the throttle firing as late as possible

and staying on maximum until the lander touches the ground, as this makes the most of the breaking force provided by the atmosphere. Conversely the minimum force on the lander would require firing as early as possible, and the fuel just running out as the lander touches down.

To tune these values to their optimum, an interval bisection algorithm was set up that calculates the value of K_h for the preferred optimisation. The function itself creates a copy of the lander object at the instance the function is called (by using the class explained in section 6.1), and running the simulation repeatedly with different values of K_h based on the interval bisection algorithm. Importantly a timeout clause is added to stop the function if the simulation time in the function exceeds a certain amount, as in scenarios such as stable orbits a crash will never be registered.

Other pitfalls in developing this calculation were noting that the lander touches down at 0.5m rather than 0 due to the height of the lander. Also ensuring that the bisection algorithm does skip the 'valid' range of landing by checking what kind of failure had occurred i.e. running out of fuel or not firing soon enough, and correcting the mid value appropriately.

This value is recalculated if a change in circumstance is forced such as removing the parachute or disabling then re-enabling the autopilot. A toggle key of 'm' was added to allow the user to select the preferred landing optimisation.

Table 1 shows tuned values for K_h when there is no parachute available and K_p is kept constant. The values for this table were determined after the introduction of planetary rotation.

Table 1 Values of K_h at K_p of 1 that provide best fuel efficiency and softest landing

Scenario	Without Parachute		With Parachute	
	Most Fuel Efficient	Softest Landing	Most Fuel Efficient	Softest Landing
1	0.03931	0.01374	0.15273	0.01030
3	0.02001	0.01735	0.05588	0.01382
4	0.03486	0.01637	0.15160	0.01128
5	0.01909	0.01670	0.05305	0.01440

5.5 AUTOPILOT MODES

The main autopilot mode can be considered Orbital Descent. There are two additional things the autopilot can do:

1. Orbital Re-Entry: bring the lander from any stable orbit into a ballistic trajectory that leads to Orbital Descent, or Orbital Transfer depending on the input radius.
2. Orbital Injection: Going from a ballistic trajectory into a stable orbit i.e. the reverse of Re-entry.

Orbital Re-Entry

From any stable circular orbit, the autopilot will request an input radius (as a factor of Mars' Radius) to which it will transfer its orbit to. The manoeuvre has 2 bursts: the first to bring the lander into an appropriate elliptical orbit with perigee or apogee at the given radius (depending on whether the lander is going to a higher orbit or not), and the second to make the orbit circular upon arrival at the

appropriate altitude. If the given radius leads to the orbiter passing within the Exosphere of Mars, then the autopilot will automatically switch to descent mode allowing it to land safely as described in 4.2.

In any stable orbit (elliptical or otherwise), the moment of momentum and the energy of the orbiter is conserved:

$$\frac{mv^2}{2} + \frac{GMm}{r} = \text{const} \quad [12]$$

$$m\mathbf{v} \times \mathbf{r} = \text{const} \quad [13]$$

For elliptical orbits, inputting the perigee and apogee points gives the following identities:

$$v_a r_a = v_p r_p \quad [14]$$

$$v_a^2 = \frac{2GM r_p}{r_a(r_a + r_p)} \quad [15]$$

In a circular orbit, the radius will be equal to either the perigee or apogee of the upcoming elliptical orbit, and the other radius is then supplied by the user, meaning that the required velocity for the orbiter to accelerate/decelerate to can be calculated. Once the descent speed reaches 0 again, the lander must be at the opposite radius and the throttle will fire once more to bring the lander to the speed required for circular orbit as shown in [2].

Because the throttle has a finite max thrust, the final circular orbit will not be perfect, as the orbiter will have moved from apogee/perigee before the burst has completed, introducing a radial component to the velocity. To correct this, the autopilot will move to the Orbital Injection algorithm to tighten up the orbit.

Orbital Injection

To get into a stable orbit from a stationary position, the lander must be accelerated to a desired altitude with no radial velocity and angular velocity as described in [2]. First of all, you can define the desired velocity as a function of altitude like in section 4.2. However, this time we have no atmosphere to work in our favour of reducing velocity, and both components of the velocity must be carefully controlled.

$$\mathbf{v}_0 = \frac{GM}{r_0} \mathbf{e}_\theta + (r_0 - r) \mathbf{K}_h \quad [16]$$

$$\boldsymbol{\varepsilon} = \mathbf{v}_0 - \mathbf{v} \quad [17]$$

$$\Delta = \left(\frac{GMm}{|\mathbf{r}|^2} \hat{\mathbf{r}} \cdot \hat{\mathbf{v}} \right) \frac{\hat{\mathbf{v}}}{T_{\max}} \quad [18]$$

Variables above with 0 subscripts are the target values we wish to achieve.

As shown above, this issue can be easily dealt with by considering the error as a vector in terms of radial and angular unit vectors. From this vector, the required angle of the lander can be determined and the output throttle regime defined.

$$\mathbf{P}_{out} = K_p \boldsymbol{\varepsilon} \quad [19]$$

$$\theta = \begin{cases} \cos^{-1}[(\widehat{\mathbf{P}_{out}} + \Delta) \cdot \mathbf{e}_r] & \boldsymbol{\varepsilon} \cdot \mathbf{e}_\theta < 0 \\ -\cos^{-1}[(\widehat{\mathbf{P}_{out}} + \Delta) \cdot \mathbf{e}_r] & \boldsymbol{\varepsilon} \cdot \mathbf{e}_\theta \geq 0 \end{cases} \quad [20]$$

$$\text{throttle} = \begin{cases} |\mathbf{P}_{out} + \Delta| & |\mathbf{P}_{out} + \Delta| < 1 \\ 1 & |\mathbf{P}_{out} + \Delta| \geq 1 \end{cases} \quad [21]$$

5.6 ENGINE LAG AND DELAY

Changing the values of 'ENGINE_LAG' and 'ENGINE_DELAY' in 'lander.h' will introduce instabilities into the system. Delay represents the time taken for a signal to reach the actuator (in this case the thrust produced in the rocket) and the lag represents how fast the actuator can respond.

Initially changing these values caused a lot of problems because of how the lander class worked, and the use of pointers which will interfere with different instances of the class. As a result, the tuner had to be modified to create a copy of the buffer and restore it at the end of the function call as well as ensure that the destructor only deletes the variable when the real object is deleted and not a virtual instance used in the tuner.

An engine lag of 5s and delay of 2s can be toggled on or off using 'c' and 'v' key respectively. The autopilot will predict up to 5s ahead to try and counteract the effects of the delay, though it will not tune due to the time it takes to run through – as it is effectively doing 100 times more calculations per time step. The delay can cause inconsistent results that will sometimes land successfully and not, this is purely down to timing of when the autopilot is activated.

6. CODING PRACTICE

As well as being an exercise in control theory, first and foremost this task is about developing an understanding of the C++ coding language. The following section details changes made to the original source-code that, while not adding extra functionality, improves readability of the code and its structure.

6.1 ORBITER AND LANDER CLASSES

Being an object-oriented programming language, classes allow groups of functions and variables to be kept in a single location, as well as control what can and cannot change those variables. All the lander variables and corresponding physics functions (including thrust and attitude stabilisation) were merged into a class that contains all the basic properties of the lander. It was then possible to remove repeated calculations of this like altitude and drag and put them in one location, making debugging more straightforward as only 1 location needs to be changed, as well as minor performance improvements by not recalculating dependant values multiple times in a loop.

All of this is put in the 'orbiter_class.h' file, though the functions where applicable were kept in their original location to maintain the code structure. Geometric arguments such as position and

velocity are kept as protected so that only the numerical dynamics function and resetting the scenario can explicitly change it to reduce the likelihood of accidentally changing values somewhere unknown.

The number of files declared in the header files could be slashed, as now only one declaration of a lander variable is required, and it will be initialised to contain all the appropriate variables.

6.2 HEADER FILES

Header files allow functions and global variables to be forward declared so that the order they are defined does not need to be considered, and that functions from different files can be accessed. Best practice suggests that every `.cpp` file should have its own `.h` file with the same name so that it is clear where declarations are kept, and do simplify any includes. The `'lander.h'` file was therefore split into 2 separate header files for the corresponding original `.cpp` files.

Now with multiple files and being included at the top of each file, repeated declarations will cause errors in compilation. Head guards using the pre-processor `#ifndef` create conditional inclusion in the file if it has not been already (from being included within a header file previously defined). It should be noted `#pragma once` can also be used but not all compilers have this pre-processor directive.

6.3 GIT

It is very easy in any coding project to make one change that breaks an entire programme. While for the core tasks only small sections of the code need to be changed, for some of the extension tasks it is very easy to ruin all the work. In cases such as this is, it is important to be able to keep track of changes and version control so that if necessary you can revert back to the last working version and try again. This was particularly crucial when creating the lander class as every file needed to be changed in some way to call the methods instead of global variables and functions, so forgetting to change one small thing can ruin the physics or stop the code compiling altogether.